

# LSTM\_Model\_3

July 21, 2020

```
[2]: import numpy as np
import pandas as pd

from tensorflow.keras.layers import (LSTM, Input, Embedding,
                                     Dense, Flatten, Concatenate,
                                     Dropout, Conv1D)
from tensorflow.keras.models import Model
from tensorflow.keras.utils import plot_model
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
[4]: from sklearn.model_selection import train_test_split

data = pd.read_csv('/content/drive/My Drive/colab resources/AppliedAI/Donors_
↳Choose/preprocessed_data.csv')

train, test = train_test_split(data, test_size=0.2, random_state=0)
```

```
[5]: word_counts = train['essay'].str.split().map(len)
np.percentile(word_counts, 97)
```

[5]: 245.0

97% of all essays have a word count of less than 250.

```
[6]: NUM_WORDS = 20_000
MAX_SEQUENCE_LENGTH = 250

tokenizer = Tokenizer(num_words=NUM_WORDS)

tokenizer.fit_on_texts(train['essay'])

train_essay_seq = tokenizer.texts_to_sequences(train['essay'])
test_essay_seq = tokenizer.texts_to_sequences(test['essay'])

train_essay_seq = pad_sequences(
    train_essay_seq, maxlen=MAX_SEQUENCE_LENGTH, dtype='uint16', padding='pre',
    ↳truncating='pre', value=0)
```

```
)
test_essay_seq = pad_sequences(
    test_essay_seq, maxlen=MAX_SEQUENCE_LENGTH, dtype='uint16', padding='pre',
    →truncating='pre', value=0
)
```

```
[8]: EMBEDDING_DIM = 50
```

```
def create_glove_embeddings(glove_path):
    embeddings_dict = {}
    with open(glove_path) as glove:
        for line in glove:
            values = line.split()
            word = values[0]
            coefs = np.asarray(values[1:], dtype='float32')
            embeddings_dict[word] = coefs

    return embeddings_dict

embeddings_dict = create_glove_embeddings(f'glove.6B.{EMBEDDING_DIM}d.txt')
```

```
[9]: def get_embedding_matrix(embeddings_dict, tokenizer):
```

```
    word_index = tokenizer.word_index
    num_words = len(word_index) + 1

    embedding_matrix = np.zeros((num_words, EMBEDDING_DIM))

    for word, index in word_index.items():
        embedding_vector = embeddings_dict.get(word)
        if embedding_vector is not None:
            embedding_matrix[index] = embedding_vector

    return embedding_matrix
```

```
embedding_matrix = get_embedding_matrix(embeddings_dict, tokenizer)
```

```
[10]: num_words = len(tokenizer.word_index) + 1
```

```
essay_inp = Input(shape=MAX_SEQUENCE_LENGTH, name='essay_inp')

x = Embedding(input_dim=num_words, # vocab size
              output_dim=EMBEDDING_DIM,
              weights=[embedding_matrix],
              trainable=False,
              input_length=MAX_SEQUENCE_LENGTH,
```

```

        name='essay_emb')(essay_inp)

x = LSTM(32, name='LSTM')(x)

essay_flat = Flatten(name='essay_flat')(x)

```

```

[11]: from sklearn.preprocessing import OneHotEncoder

one_hot = OneHotEncoder(sparse=False, handle_unknown='ignore')

train_cats = train.select_dtypes(object).drop('essay', axis=1)
test_cats = test.select_dtypes(object).drop('essay', axis=1)

one_hot.fit(train_cats)

train_ohe = one_hot.transform(train_cats)
test_ohe = one_hot.transform(test_cats)

```

```

[12]: train_nums = train.select_dtypes(exclude=object).drop('project_is_approved',
    ↪axis=1)
test_nums = test.select_dtypes(exclude=object).drop('project_is_approved',
    ↪axis=1)

train_non_txt = np.c_[train_ohe, train_nums]
test_non_txt = np.c_[test_ohe, test_nums]

```

```

[13]: train_non_txt = np.expand_dims(train_non_txt, 2)
test_non_txt = np.expand_dims(test_non_txt, 2)

input_dim = train_non_txt.shape[1:]
non_txt_inp = Input(shape=input_dim, name='non_txt_inp')
x = Conv1D(128, 3, activation='relu', name='first_conv1d')(non_txt_inp)
x = Conv1D(64, 5, activation='relu', name='second_conv1d')(x)
non_txt_flat = Flatten(name='non_txt_flat')(x)

```

```

[14]: x = Concatenate(name='concat')([essay_flat, non_txt_flat])
x = Dense(64, activation='relu', name='dense_1')(x)
x = Dropout(0.33, name='dropout_1')(x)
x = Dense(64, activation='relu', name='dense_2')(x)
x = Dropout(0.33, name='dropout_2')(x)
x = Dense(32, activation='relu', name='dense_3')(x)
output_layer = Dense(1, activation='sigmoid', name='output_layer')(x)

```

```

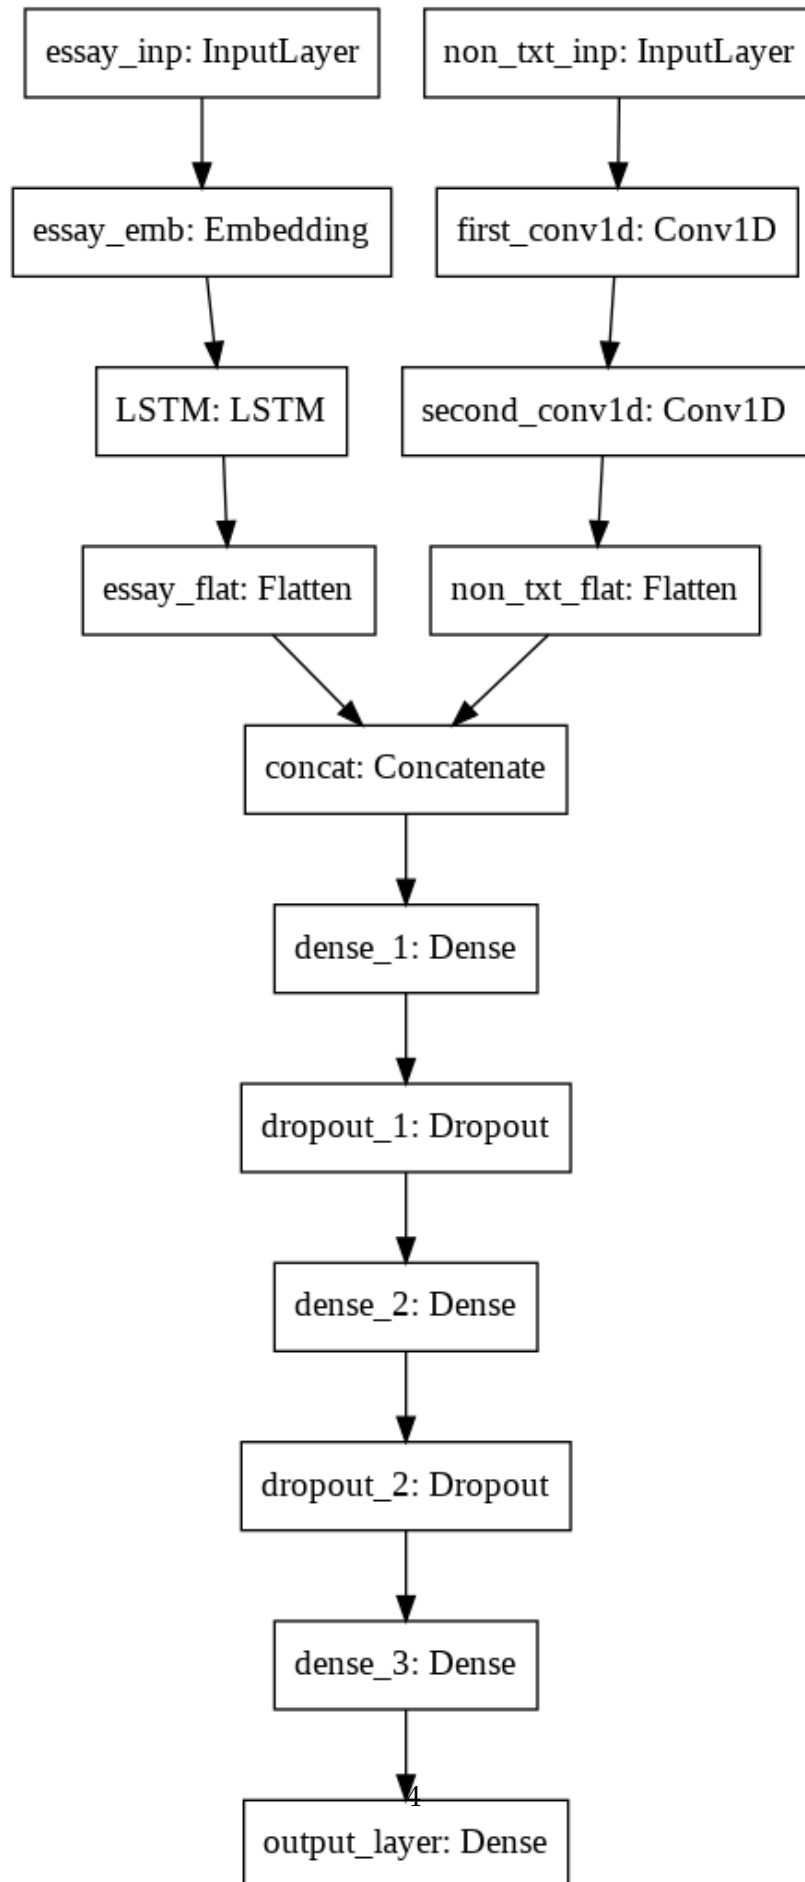
[15]: model = Model(inputs=[essay_inp, non_txt_inp], outputs=[output_layer])
model.compile(loss='binary_crossentropy', optimizer='adam')
plot_model(model)

```

```

[15]:

```



```
[16]: from sklearn.metrics import roc_auc_score
from tensorflow.keras.callbacks import (Callback, EarlyStopping,
                                         TensorBoard, ReduceLROnPlateau)
```

```
class ROCCallback(Callback):

    def __init__(self, validation_data):
        super(ROCCallback, self).__init__()
        self.validation_data = validation_data

    def on_epoch_end(self, epoch, logs={}):
        probs = self.model.predict(self.validation_data[0])
        y_true = self.validation_data[1]
        score = roc_auc_score(y_true, probs)
        logs['auc'] = score
```

```
[17]: X_train = [train_essay_seq, train_non_txt]
y_train = train['project_is_approved']

X_test = [test_essay_seq, test_non_txt]
y_test = test['project_is_approved']
```

```
[18]: roc_callback = ROCCallback((X_test, y_test))
early_stopping = EarlyStopping(patience=3)
tensorboard = TensorBoard()
reduce_lr = ReduceLROnPlateau()

callbacks = [
    roc_callback,
    early_stopping,
    tensorboard,
    reduce_lr,
]
```

```
[19]: !rm -rf ./logs/*

history = model.fit(X_train, y_train, batch_size=32, epochs=100,
                    callbacks=callbacks, validation_data=(X_test, y_test))
```

```
Epoch 1/100
2732/2732 [=====] - 155s 57ms/step - loss: 0.4069 -
val_loss: 0.3904 - auc: 0.7247 - lr: 0.0010
Epoch 2/100
2732/2732 [=====] - 155s 57ms/step - loss: 0.3825 -
val_loss: 0.3773 - auc: 0.7399 - lr: 0.0010
```

Epoch 3/100  
2732/2732 [=====] - 155s 57ms/step - loss: 0.3753 -  
val\_loss: 0.3830 - auc: 0.7482 - lr: 0.0010  
Epoch 4/100  
2732/2732 [=====] - 154s 56ms/step - loss: 0.3702 -  
val\_loss: 0.3694 - auc: 0.7539 - lr: 0.0010  
Epoch 5/100  
2732/2732 [=====] - 153s 56ms/step - loss: 0.3651 -  
val\_loss: 0.3732 - auc: 0.7516 - lr: 0.0010  
Epoch 6/100  
2732/2732 [=====] - 154s 56ms/step - loss: 0.3611 -  
val\_loss: 0.3679 - auc: 0.7573 - lr: 0.0010  
Epoch 7/100  
2732/2732 [=====] - 153s 56ms/step - loss: 0.3580 -  
val\_loss: 0.3756 - auc: 0.7616 - lr: 0.0010  
Epoch 8/100  
2732/2732 [=====] - 154s 56ms/step - loss: 0.3537 -  
val\_loss: 0.3662 - auc: 0.7623 - lr: 0.0010  
Epoch 9/100  
2732/2732 [=====] - 153s 56ms/step - loss: 0.3501 -  
val\_loss: 0.3697 - auc: 0.7592 - lr: 0.0010  
Epoch 10/100  
2732/2732 [=====] - 154s 56ms/step - loss: 0.3473 -  
val\_loss: 0.3757 - auc: 0.7580 - lr: 0.0010  
Epoch 11/100  
2732/2732 [=====] - 153s 56ms/step - loss: 0.3433 -  
val\_loss: 0.3752 - auc: 0.7542 - lr: 0.0010